



# Lecture 3: Deep Learning Software and PyTorch tutorial on Deep Learning 2023

**Marko Savic, Zhuo Su and Jukka Komulainen**

06 November 2023





Deep Learning Software

Basic concepts for machine learning

Coding examples

Resources

Assignment 1





# Deep learning frameworks



- High-level programming interfaces
  - *Python/NumPy instead of low-level programming with C, C++, CUDA or HIP (CPU + GPU)*
- Provide seamless CPU / GPU usage
  - *Multi-GPU & distributed training of deep neural networks*
- Open source
  - *Very active communities*





# Deep learning frameworks



- In fact, tools for defining **static** or **dynamic** general-purpose **computational graphs**
- **Automatic differentiation!**
  - *no need to compute the partial derivatives of millions of parameters by hand*
- All you need to do is to implement forward propagation through a computational graph and specify the cost function and optimizer you want to use!
- The DL frameworks will then compute the derivatives for you, by moving backwards through the graph





# Several different options



- Different companies and institutes have created their own:
  - *PyTorch*
  - *TensorFlow*
  - *Theano*
  - *Jax*
  - *Caffe*
  - *MXNet*
  - *MATLAB*
  - ...
- PyTorch and TensorFlow are the most popular ones





vs



TensorFlow




- **PyTorch (by Facebook)**
  - *Only low-level API: easier than TensorFlow but more low-level than keras*
  - *Easy to debug*
  - *Allows more control and customization, easier experimentation with new architectures*
  - *Great for research use & rapid prototype development*
  - *Catching up with features that are already mature in TensorFlow*
- **TensorFlow (by Google)**
  - *Low-level API but Keras provides a nice high-level interface*
  - *Difficult to debug*
  - *Keras is easier than PyTorch if you just want to apply deep learning, and not do research in machine learning*
  - *Widely used at production level in industry*
  - *For large-scale deployment*
  - *Lightweight models via TensorFlow Lite*




## Google Search Trends PyTorch vs TensorFlow

Pytorch: (Worldwide) TensorFlow: (Worldwide)





## How to implement a deep learning algorithm



We learn how to solve a problem from the data. Specifically, we train a model (using a training scheme) to fit the data






Problem: Which handwritten digit is in this image?



input



model



0

output




[illegible]




Problem: Which handwritten digit is in this image?

Data: MNIST database of handwritten digits



<http://yann.lecun.com/exdb/mnist/>

Model  
(randomly initialized)



Training scheme



Learning rate?  
Optimizer (next lectures)?  
Number of iteration?  
Loss function?





Before we dig it deeper, let's see an easier case first...





Data: 5 points

x	y	value
0.8345	0.9785	8.4596
0.0993	0.6754	2.2981
1.8054	1.8001	13.8385
1.8896	0.7385	11.3696
0.9817	0.2224	4.7279

Model: a plane  $f(x, y) = ax + by + c$





Data: 5 points

x	y	value (ground truth)
0.8345	0.9785	8.4596
0.0993	0.6754	2.2981
1.8054	1.8001	13.8385
1.8896	0.7385	11.3696
0.9817	0.2224	4.7279

Model: a plane

$$f(x, y) = ax + by + c$$



Fitting error:

For point 1:  $e_1 = |v_1 - f(x_1, y_1)| = |v_1 - (ax_1 + by_1 + c)|$

For point i:  $e_i = |v_i - f(x_i, y_i)| = |v_i - (ax_i + by_i + c)|$

Total error:  $cost = \frac{1}{N} \sum_{i=1}^N e_i = \frac{1}{N} \sum_{i=1}^N |v_i - (ax_i + by_i + c)|$

The smaller the cost, the better the model

How can we find the parameters a, b, c to minimize the cost?





The solution is simple: using gradients



$$g_a = \frac{\partial \text{cost}}{\partial a} = \frac{1}{N} \sum_{i=1}^N \frac{\partial e_i}{\partial a} = \frac{1}{N} \sum_{i=1}^N \frac{\partial |v_i - (ax_i + by_i + c)|}{\partial a}$$

$$g_b = \dots$$

$$g_c = \dots$$

Gradient descent algorithm (updating rule for the coefficients):

*Define cost, learningrate  $\eta$  and number of iterations  $n\_iter$*

*$a = \text{rand}()$*

*$b = \text{rand}()$*

*$c = \text{rand}()$*

*for  $i$  in range( $n\_iter$ ):*

*calculate  $g_a, g_b, g_c$*

*$a = a - \eta g_a$*

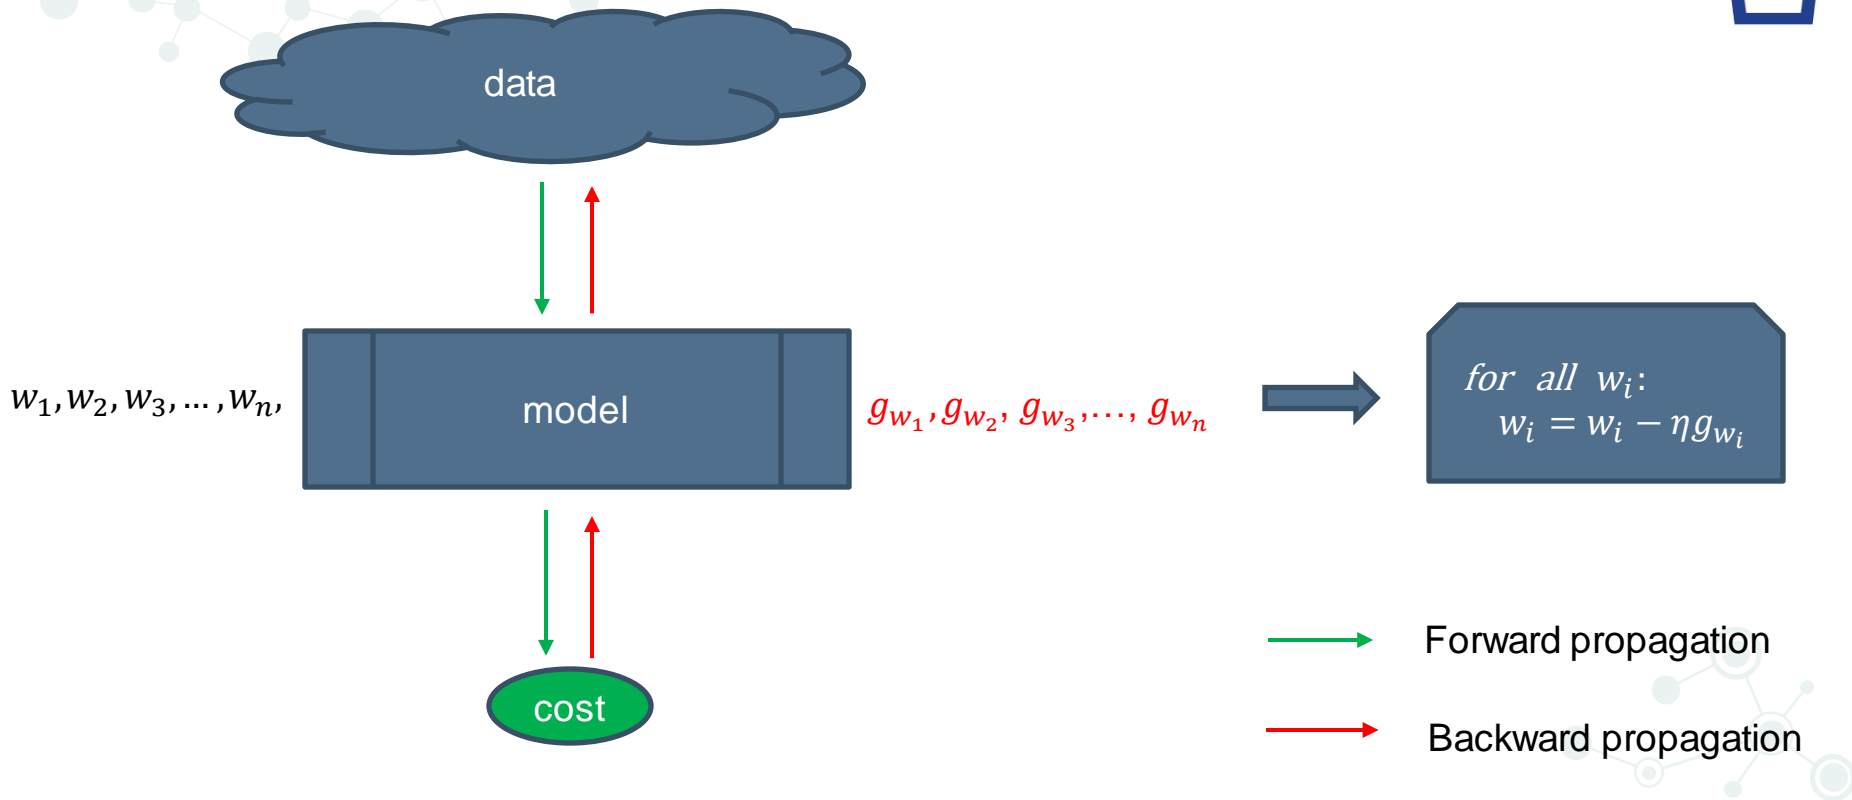
*$b = b - \eta g_b$*

*$c = c - \eta g_c$*

*return  $a, b$  and  $c$*



# Gradient descent algorithm in PyTorch program:






# What are tensors?



- Generalization of matrices to  $n$  dimensions
  - *1D tensor: vector*
  - *2D tensor: matrix*
  - *3D, 4D, 5D tensors*
- A set of  $N$ ,  $C$  channel  $H \times W$  images can be a  $[N, C, H, W]$  4D tensor





A simple example of code in PyTorch  
(please find the script shared in Moodle after the lecture).


When the TA goes on  
about boring math  
stuff, but still hasn't  
taught you how to  
generate anime  
waifus/husbandos






Our targeted problem: Which number did you write?

Data: MNIST database of handwritten digits



<http://yann.lecun.com/exdb/mnist/>

Model  
(randomly initialized)



Training scheme

Learning rate?  
Optimizer?  
Number of iteration?  
Loss function?





input



100 x 1 x 28 x 28



$w_{11}, w_{12}, w_{13}, \dots, w_{1n}$

Conv  
layer 1

100 x 24 x 14 x 14

$w_{21}, w_{22}, w_{23}, \dots, w_{2p}$

Conv  
layer 2

100 x 32 x 7 x 7

$w_{31}, w_{32}, w_{33}, \dots, w_{3q}$

FC layer

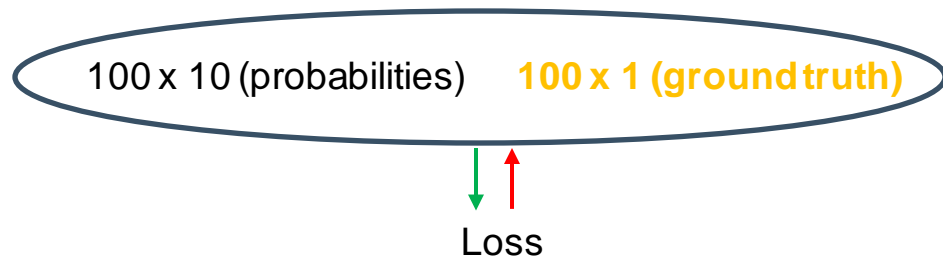
100 x 10 (logits)

Softmax

100 x 10 (probabilities)    100 x 1 (ground truth)



Loss





0 1 2 3 4 5 6 7 8 9  
[0.1, 0.05, 0.6, 0.05, 0.07, 0.03, 0.02, 0.02, 0.05, 0.01]

0 1 2 3 4 5 6 7 8 9  
[0, 0, 1.0, 0, 0, 0, 0, 0, 0, 0]



Cross entropy (or KL divergence)



Now, let's take a look at the code  
(please find the script shared in Moodle after the lecture).





# Useful tips:



60-minute tutorial:

[https://pytorch.org/tutorials/beginner/deep\\_learning\\_60min\\_blitz.html](https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html)

Use Google anytime you come across problems

Attend the lectures, most concepts and fundamentals will be covered

If you have any questions or need help, write on discord or contact us by email





# What resources you can use:



- You own PC (next slide contains details about setting up an environment)
- GPU PCs in TS135 (when vacant, or during our booked slot in December)
- Google Colab with University google account
- CSC notebooks





# Environment DL23 and Jupyter



Creating the environment:

Step 1: Go to anaconda website and install miniconda

<https://docs.conda.io/projects/miniconda/en/latest/>

Step 2: Create isolated conda environment (use supplied .yaml file from moodle)

```
conda env create -f environment.yaml  
source activate DL23
```

OR

In Anaconda Navigator go to  
Environments -> Import

Using it:

From command line (or also simply from UI if you have windows):

```
conda activate DL23
```

After that run jupyter notebook from the activated DL23 environment:

```
jupyter-notebook
```

OR

```
jupyter-lab
```





# Cloud resources



Google Colab:

Step 1: If you haven't already, follow patio instructions on how to get university google account:  
<https://patio.oulu.fi/en/services-and-instructions/it-services/information-systems/google-workspace-education>

Step 2: Go to:  
<https://colab.research.google.com>

Step 3: Upload the notebook you want to work on.

CSC Notebooks:

Step 1: Go to and login:  
<https://notebooks.rahtiapp.fi/>

Step 2: Select "Practical Machine learning" and run

Step 3: Upload the notebook you want to work on.





Let's take a quick look at assignment 1





Thanks for listening

